

COMP9727: Recommender Systems

Lecture 5: Social Recommender Systems

Wayne Wobcke

e-mail: w.wobcke@unsw.edu.au

This Lecture

Recommendation using social media data, often on social media

- Social Filtering, using real social networks (not “neighbourhoods”)
 - ▶ Facebook, Instagram, YouTube, Twitter/X, Reddit, Quora, ...
 - ▶ “Web 2.0” = “User Generated Content” = “Wisdom of the Crowd”
 - ▶ Use “likes”, retweets, upvotes, tags, comments, reviews, ...
- Applications
 - ▶ Feeds, News, Blogs, Photos, Videos, Restaurants, Q&A, ...
 - ▶ Tag-aware recommendation (from a “folksonomy”)
- Research Problems
 - ▶ Networks: “Trust” (Reputation) Propagation (c.f. Amazon)
 - ▶ NLP: Sentiment Analysis, Stance Detection

YouTube Recommender (2010)

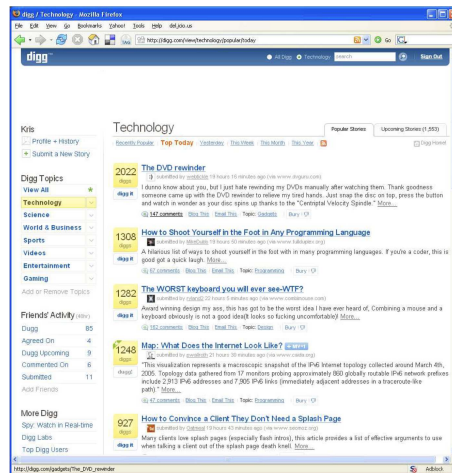
1. Use co-visitation (association rule [lift](#)) in last 24 hours (offline)
 - Seed video gives up to N related videos with score > threshold
2. Get [seed set](#) = watched/liked videos from user activity
3. Generate [candidate set](#) by applying co-visitation
4. Expand a few more times to get more “diverse” candidates
5. Rank candidates using linear function for this user
 - Quality: views, ratings, comments, favourites, shares, upload time
 - Specificity: views, watch time of seed video in user history
6. Diversity: Remove videos in candidate set too similar to each other
 - Same seed video, same uploader, same content, etc.

Digg News Aggregator

- Can “follow” other Digg users: follower/followee (like Twitter)
 - ▶ MySpace (music), Flickr (photos), del.icio.us (web bookmarks)
 - ▶ [Quiz Question](#): How many of these still exist?
- Has same “front page” for all users
- Rank users by number of stories on front page
- Rank stories by number of diggs and rate of diggs, etc.
 - ▶ [Open to manipulation by coordinated teams \(extremists\)](#)
- Allow users to view articles based on friends’ activity
 - ▶ In turn boosting those articles to the front page ...

Similar to Reddit: feed not personalized

Digg Front Page



UNSW

©W. Wobcke et al. 2023–2025

Trust Propagation

- Suppose values are between 0 and 1
- Multiply values along a path such as $A \rightarrow B \rightarrow C$
 - Or discount $B \rightarrow C$ by some decay factor $\beta \lesssim 1$
- Aggregate values across multiple paths
 - Such as max, min, weighted average, etc., with or without decay
 - With multiple paths between two nodes, choose all shortest ones?
- Application
 - Use trust rather than similarity in user-based CF
 - Determine neighbourhoods using trust in item-based CF
 - Can have a different trust network for each topic (domain experts!)

UNSW

©W. Wobcke et al. 2023–2025

Trust Networks

- How much do you “trust” your “friends”?
- Network with arrow $A \rightarrow B$ as A trusts B (to some degree)
- Trust propagation
 - e.g. transitivity: $A \rightarrow B$ and $B \rightarrow C$ implies $A \rightarrow C$
- Questions
 - What is the role of social media “influencers”?
 - Who influences the influencers?
 - Lazarsfeld & Katz (1955) *Personal Influence*: opinion leaders


Quiz Question: What is the #1 predictor of whether you own an iPhone?

Quiz Question: Is influence domain specific or general?

UNSW

©W. Wobcke et al. 2023–2025

MoleTrust



UNSW

©W. Wobcke et al. 2023–2025

Google PageRank

- Importance “flows” from one node to another
- Importance of a node derived from in-coming links
- $rank(j) = \sum_i \frac{rank(i)}{d_i}$ over nodes $i \rightarrow j$ where d_i is out-degree of i
- In matrix form, solution is π such that $P\pi = \pi$ (eigenvector!)
- Solve by repeated powers $P\pi, P^2\pi, \dots$ starting with random π
- Add damping factor $\beta \lesssim 1$ to avoid “dead ends”


Tag-Aware Recommendation

Folksonomy = Collection of tags (over a set of items)


- Recommend items to user based on tagging behaviour
- Recommend items to user for a given tag
- Recommend tags to user based on searching
- Recommend users to users based on similar tags used
- Recommend tags to user for a given item

Use content-based recommendation and/or CF – for which scenarios?

Problem: Folksonomies are large, complex and messy



PageRank Example



$$P = \begin{pmatrix} 0 & 0 & 0 & 0 & \frac{1}{3} \\ 0 & 0 & \frac{1}{2} & 1 & \frac{1}{3} \\ \frac{1}{3} & 0 & 0 & 0 & \frac{1}{3} \\ \frac{1}{3} & 0 & \frac{1}{2} & 0 & 0 \\ \frac{1}{3} & 1 & 0 & 0 & 0 \end{pmatrix} \quad R = \begin{pmatrix} \frac{1-\beta}{5} & \frac{1-\beta}{5} & \frac{1-\beta}{5} & \frac{1-\beta}{5} & \frac{1}{3}\beta + \frac{1-\beta}{5} \\ \frac{1-\beta}{5} & \frac{1-\beta}{5} & \frac{1}{2}\beta + \frac{1-\beta}{5} & \beta + \frac{1-\beta}{5} & \frac{1}{3}\beta + \frac{1-\beta}{5} \\ \frac{1}{3}\beta + \frac{1-\beta}{5} & \frac{1-\beta}{5} & \frac{1-\beta}{5} & \frac{1-\beta}{5} & \frac{1}{3}\beta + \frac{1-\beta}{5} \\ \frac{1}{3}\beta + \frac{1-\beta}{5} & \frac{1-\beta}{5} & \frac{1}{2}\beta + \frac{1-\beta}{5} & \frac{1-\beta}{5} & \frac{1-\beta}{5} \\ \frac{1}{3}\beta + \frac{1-\beta}{5} & \beta + \frac{1-\beta}{5} & \frac{1-\beta}{5} & \frac{1-\beta}{5} & \frac{1-\beta}{5} \end{pmatrix}$$


e.g. $v_1 = \frac{1}{2}v_2 + v_3 + \frac{1}{3}v_4$

Tag-Aware Recommendation Example

Content-based recommendation applied to movie ratings

- Movies are documents – **tag clouds**
- User profiles are tag clouds for each rating value
- Match user profiles with tf-idf vectors for movies
- Can use weighted sum by similarity over rating values
- Can also cluster tags using similarity of tf-idf vectors
 - ... to improve recall and maybe address cold start problem

Tag Cube



Sentiment Analysis

Determine sentiment of **language**: positive/neutral/negative

- Hand-curated lexicons (accurate but incomplete)
- Topic modelling (more complete)
 - ▶ Topics are probability distributions over words
 - ▶ Need good “seed words” to define topics
- Neural networks with Word2Vec and Multi-Layer Perceptron
 - ▶ Word2Vec is “self supervised”; training data for sentiment
 - ▶ Example in tutorials uses only positive/negative (easier!)

Tag-Aware Recommendation Examples

CF applied to movie ratings and tags (c.f. hybrid methods)

- With limited data ... (here 0/1 matrix)
 - ▶ Add tags into the ratings matrix as if they are items (**pseudo-items**)
 - ▶ Add tags into the ratings matrix as if they are users (**pseudo-users**)
 - ▶ Use weighted sum of user-based and item-based predictions
 - ▶ c.f. feature augmentation in hybrid recommender systems
- With more data ...
 - ▶ Apply “query expansion” by tag clusters to user’s tags
 - ▶ Add expanded tags into the ratings matrix for each movie
 - ▶ Augment user similarity with tf-idf similarity on tags
 - ▶ c.f. meta-features in hybrid recommender systems

Artificial Neural Networks

(Artificial) Neural Networks are made up of nodes which have

- Input edges, each with some **weight**
- Output edges (with **weights**)
- An **activation level** (a function of the inputs)

Weights can be positive or negative and may change over time (learning)

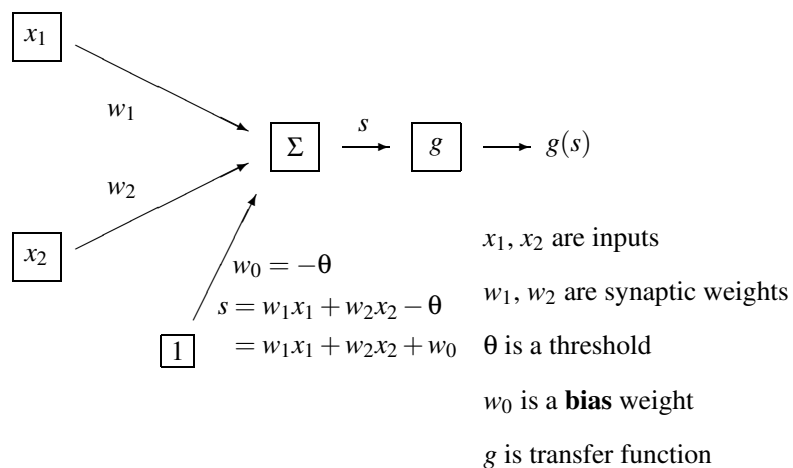
The **input function** is the weighted sum of the activation levels of inputs

The activation level is a non-linear **transfer function** g of this input


$$\text{activation}_i = g(s_i) = g\left(\sum_j w_{ij}x_j\right)$$

Some nodes are inputs (sensing), some are outputs (action)


McCulloch & Pitts Model of a Single Neuron



Simple Neural Network



Transfer Functions



Note: if $g(z) = \frac{1}{1 + e^{-z}}$ $g'(z) = g(z)(1 - g(z))$

$g(z) = \max(z, 0)$ $g'(z) = \begin{cases} 1 & \text{if } z \geq 0 \\ 0 & \text{if } z < 0 \end{cases}$

Gradient Descent

Define an **error** (or “loss”) function E as (half) the sum over all input patterns of the square of the difference between actual output and desired output

$$E = \frac{1}{2} \sum (z - t)^2$$


The aim is to find a set of weights for which E is very low.

If the functions are smooth, use multi-variate calculus to define how to adjust the weights so error moves in steepest downhill direction

$$w \leftarrow w - \eta \frac{\partial E}{\partial w}$$


Parameter η is called the **learning rate**

Local Search in Weight Space



Problem: Because of the step function, the landscape will not be smooth but will instead consist almost entirely of flat local regions and “shoulders”, with occasional discontinuous jumps

Forward Pass



$$\begin{aligned} u_1 &= b_1 + w_{11}x_1 + w_{12}x_2 \\ y_1 &= g(u_1) \\ s &= c + v_1y_1 + v_2y_2 \\ z &= g(s) \\ E &= \frac{1}{2} \sum (z - t)^2 \end{aligned}$$

Chain Rule

If

$$y = y(u)$$

$$u = u(x)$$

Then

$$\frac{\partial y}{\partial x} = \frac{\partial y}{\partial u} \frac{\partial u}{\partial x}$$

This principle can be used to compute the partial derivatives in an efficient and localized manner. The transfer function must be differentiable.

$$\text{Note: if } z(s) = \frac{1}{1 + e^{-s}} \quad \frac{\partial z}{\partial s} = z(1 - z)$$

$$\text{if } z(s) = \tanh(s) \quad \frac{\partial z}{\partial s} = 1 - z^2$$

Backpropagation

Partial Derivatives

$$\frac{\partial E}{\partial z} = z - t$$

$$\frac{dz}{ds} = g'(s) = z(1 - z)$$

$$\frac{\partial s}{\partial y_1} = v_1$$

$$\frac{dy_1}{du_1} = y_1(1 - y_1)$$

Useful notation

$$\delta_{\text{out}} = \frac{\partial E}{\partial s} \quad \delta_1 = \frac{\partial E}{\partial u_1} \quad \delta_2 = \frac{\partial E}{\partial u_2}$$

Then

$$\delta_{\text{out}} = (z - t) z (1 - z)$$

$$\frac{\partial E}{\partial v_1} = \delta_{\text{out}} y_1$$

$$\delta_1 = \delta_{\text{out}} v_1 y_1 (1 - y_1)$$

$$\frac{\partial E}{\partial w_{11}} = \delta_1 x_1$$

Partial derivatives can be calculated efficiently by backpropagating deltas through the network

Cross Entropy

For classification tasks, target t is either 0 or 1, so better to use

$$E = -t \log(z) - (1-t) \log(1-z)$$

This can be justified mathematically, and works well in practice – especially when negative examples vastly outweigh positive ones.

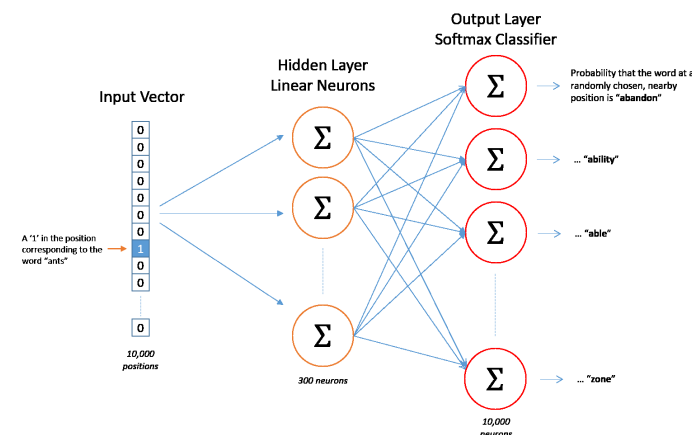
It also makes the backpropagation computations simpler:

$$\frac{\partial E}{\partial z} = \frac{z-t}{z(1-z)}$$

If $z = \frac{1}{1+e^{-s}}$

$$\frac{\partial E}{\partial s} = \frac{\partial E}{\partial z} \frac{\partial z}{\partial s} = z-t$$

Word2Vec Architecture

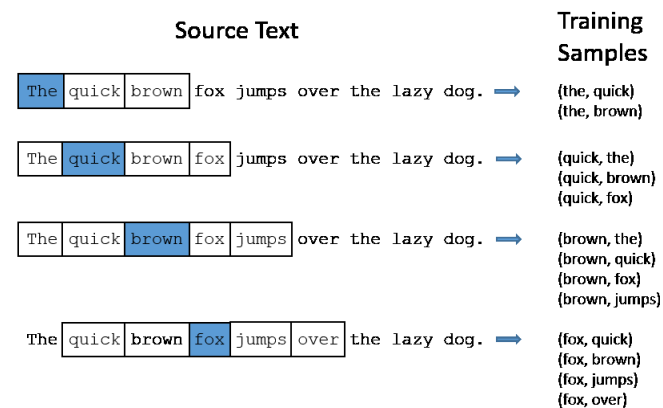


Word2Vec

- Word “embeddings” as vectors (≈ 300 dimensions)
- Training: In effect, inputs and outputs are single words
 - CBOW: Context words (words in window) predict word
 - Skip-gram: Word predicts context words (words in window)
 - Add negative training examples
- Embeddings are first layer of the network, ignore second layer
- Words with “similar” contexts have “similar” embeddings
- Alternative uses word and context word vectors (wrong?)
- Can use pretrained model (Google News) or train new model

Somewhat similar to matrix factorization with 300 dimensions

Word2Vec Training Data




Word2Vec Sentiment Embeddings

What we want to see:



Multi-Layer Perceptron (MLP)



Sentiment Analysis with Word2Vec

- Word Embeddings
 - ▶ Word vectors trained on 25,000 movie reviews
 - ▶ Concatenate word vectors for each word in the review
- Multi-Layer Perceptron
 - ▶ First layer is embedding layer
 - ▶ Hidden layer fully connected (128 dimensions)
 - ▶ Output layer fully connected (1 neuron) for sentiment (+/−)
 - ▶ Train through backpropagation
 - ▶ 6,400,257 parameters (weights) to learn

Further Research Problems

- Aspect-Based Sentiment Analysis
 - ▶ Cameras: zoom, battery, resolution, ease of use, etc.
 - ▶ Hotels: location, room, cleanliness, host, value, etc.
 - ▶ Useful for more precise recommendations, or critiquing
- Stance Detection
 - ▶ Abortion, gun control, animal rights – **for** or **against**
 - ▶ Useful for political advertising

Summary

- Suitable for social media companies
- Much money to be made from targeted advertising
- Role of “influencers” important
- Research shows using tags improves recommendation
 - ▶ ... despite folksonomies being very noisy
- Susceptible to attacks by bots, extremists
 - ▶ And echo chambers leading to polarization
- Next lecture social network recommendation